

Day 2 Exercises

Welcome to Day 2 of the LLM and Agent Development course. Today's practical exercises will help you apply the concepts we've explored and deepen your understanding through hands-on implementation. These exercises are designed to challenge your skills while providing a structured framework for success.



Practice Core Concepts

Apply theoretical knowledge to real-world scenarios through guided exercises



Build Practical Skills

Develop hands-on experience with LLM integration and agent architectures



Demonstrate Mastery


Showcase your understanding through deliverable outputs and implementations

Exercise 1: MCP Server

MCP Server with one tool: search

Based on `demo_code/mcp/basic.py`

- Keep only one tool: call the tool **search**
- Tool with two parameters: **query** and **num_results**
- Return some dummy results for now
- Remove `run_client_example()`
- In `main()` keep only the `mcp.run()` code but ensure server uses `streamable-http`
- Run the server
- Use MCP Protocol Inspector to test the MCP server (run **`npx @modelcontextprotocol/inspector`**; requires Node.js)

 See `streamable.py` for information required by the `run()` method

If you want to add authentication via a HTTP header, see `search_agent/mcp_server` for an example: adds Starlette middleware to the FastMCP server.

Azure AI Search Prerequisites


- Azure AI Search instance accessible by all users (do not use **Free Tier**)
- Azure OpenAI resource with **text-embedding-3-large** embedding model
- Storage Account accessible by all users

Exercise 2: Azure AI Search

Create an index

We will use the **Import data (new)** wizard

- Create a container in the storage account with a unique name
- Upload a number of PDFs or DOCX files to the container
- In the shared Azure AI Search instance, use the **Import data (new)** wizard with **Azure Blob Storage** as the data source and follow the steps
 - You will need access to the Azure OpenAI resource and select the embedding model
- When the wizard asks for an index name, use a **unique name** on the AI Search instance

 This should create an indexer, data source, skillset and index. Use **Search Explorer** with your index to check for documents.

Note: the skillset use a split skill and embedding skill to convert documents, split them and create vectors for each split or chunk

Exercise 3: Modify the MCP Server

Modify the search tool to search your index

Use `aisearch/query_examples.py` for sample search code

- Search your index using hybrid search + semantic reranking
- You can hard code the index name, search keys etc... in the code
- Optional: filter the query based on semantic reranker score
- Test the tool in MCP Inspector

Exercise 4: Create a search agent

Create a search agent that uses the MCP tool

See `search_agent` folder for inspiration

- You can choose how you create the agent:
 - Responses API:
 - Built-in mcp functionality: requires exposing your MCP server on the Internet (e.g. ngrok)
 - Use a function tool: custom function can use FastMCP MCP client to run the tool on the local MCP server
 - Azure AI Foundry Agent:
 - Built-in mcp functionality: similar to Responses API; requires public endpoint
 - Use a function tool: similar to Responses API
 - LangChain Agent (or other frameworks)
 - MCP functionality is not implemented at the API level
 - This means that your server **does not** require a public endpoint
 - See `search_agent/langchain_agent/search_agent_mcp.py`: uses native support for MCP in LangChain

 LangChain recommended